

- Sockets are *one* way of doing inter-process communication (IPC) between processes on the *same* or on *different* machines in Unix
- First introduced with BSD 4.1c in 1983
 - Now available in all flavors of Unix and also on Windows („winsock“)
 - Especially well suited for and often used in client-server applications
 - http, ftp, ssh, ... are sockets-based

- **Socket:** *one* endpoint of a bi-directional communication link
 - A communication link always consists of two sockets
 - Setting up a communication link: create two sockets and connect them

- **Client-server applications**
 - Non-symmetric roles
 - Server waits for a connection from a client
 - One server can usually handle many clients: `fork()`
 - Loose coupling: client and server only need to speak the same protocol, no need to use the same programming language or the same architecture on both parts

- A socket is accessed through a descriptor (system-provided integer)
 - Descriptor returned from or passed to functions
 - Many similarities with files (file descriptors); example: read and write functions
 - Difference to files in creation and control of options (more complex)
- Socket creation and usage in C: (all in `<sys/socket.h>`)
 - `socket()`
 - `socketpair()`
 - `bind()`
 - `listen()`
 - `accept()`
 - `connect()`
 - `read()`
 - `write()`
 - `close()`

- **int socket (int domain, int type, int protocol)**
 - Domain is AF_UNIX or AF_INET
 - AF_INET for processes on different or on the same host
 - AF_UNIX for processes on the same machine only
 - Type is the style of communication:
 - SOCK_STREAM: connection-oriented „stream“: no record boundaries, guaranteed delivery
 - SOCK_DGRAM: connection-less „datagram“: sending individual records, delivery is not guaranteed
 - Protocol allows the specification of the underlying protocol to be used;
 - AF_INET and stream goes with TCP
 - AF_INET and datagram goes with UDP
 - The system will choose the most appropriate when specifying 0
 - Our application uses AF_INET and SOCK_STREAM

Assing a name to a socket: bind

- `int bind(int sockfd, struct sockaddr *my_addr, socklen_t, addrlen)`
 - Assigns a name to a socket, so that it can be referenced by another process
 - In the AF_UNIX domain, a name is established using the file system
 - In the AF_INET domain a name consists of an internet (IP) address and a port number
 - Example: 132.159.32.1:80
 - Ports below 1024 are reserved

Connect to another socket: connect

- `int connect(int sockfd, const struct sockaddr *serv_addr, socklen_t addrlen);`
 - Specify the local socket (`sockfd`) and the name of the remote socket
 - Remote socket must have been bound to that name
 - When `connect` succeeds, the connection is set up and you can read/write to the socket specified by `sockfd`

Setting up a Server to accept connections


- Create (socket()) and assign name (bind())
- Change to listening mode
 - `int listen(int s, int backlog)`
 - backlog is the length of the queue where unhandled connection requests are placed
 - non-blocking
- Wait for connections
 - `int accept (int s, struct sockaddr *addr, socklen_t *addrlen);`
 - blocking
 - returns a *new* socket used to talk to client, client's address information is placed in name
 - The old socket is used to handle new connection requests

- `ssize_t read (int fd, void *buf, size_t count);`
- `ssize_t write (int fd, const void *buf, size_t count);`

- `int send (int s, const void *msg, size_t len, int flags);`
- `int recv (int s, void *buf, size_t len, int flags);`

- **Also:** `int recv(int s, void *buf, size_t len, int flags);`
 - `Recvfrom, recvmsg, sendto sendmsg...`

Client and Server



- Easier, Java hides a lot of details
- Two classes in java.net.*
 - Socket
 - ServerSocket
 - Example:

```
mysock = new Socket(<hostname>, <port>);  
servsock = new ServerSocket(<port>);
```
 - Read from and write to socket via input and output streams

```
mysock.getOutputStream()  
mysock.getInputStream()
```
- Server implementation in Java
 - Blocking accept();

```
Socket clientSocket = null;  
try {  
    clientSocket = serverSocket.accept();  
}  
catch (IOException e) {}
```